

Scaling GraphQL for 500,000,000 req/min

whoami

Tushar Mathur

- GraphQL Enthusiast
- Open Source Contributor
- Passionate about DevEx

Part 1

GraphQL Journey

2016 Dream11

- Fantasy Gaming Platform
- Early Stage
- Monolith to Microservices Considerable Scale

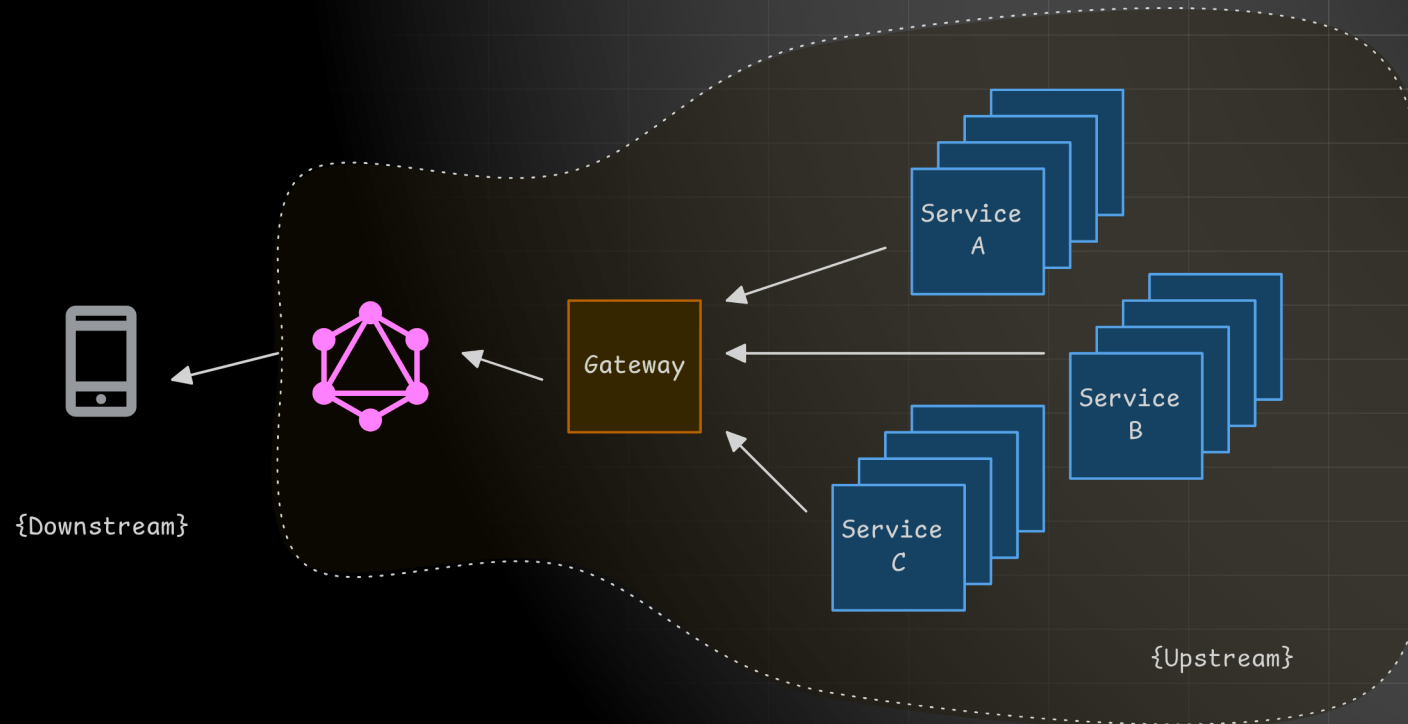
2016 Dream11

- Fantasy Gaming Platform
- Early Stage
- Monolith to Microservices Considerable Scale

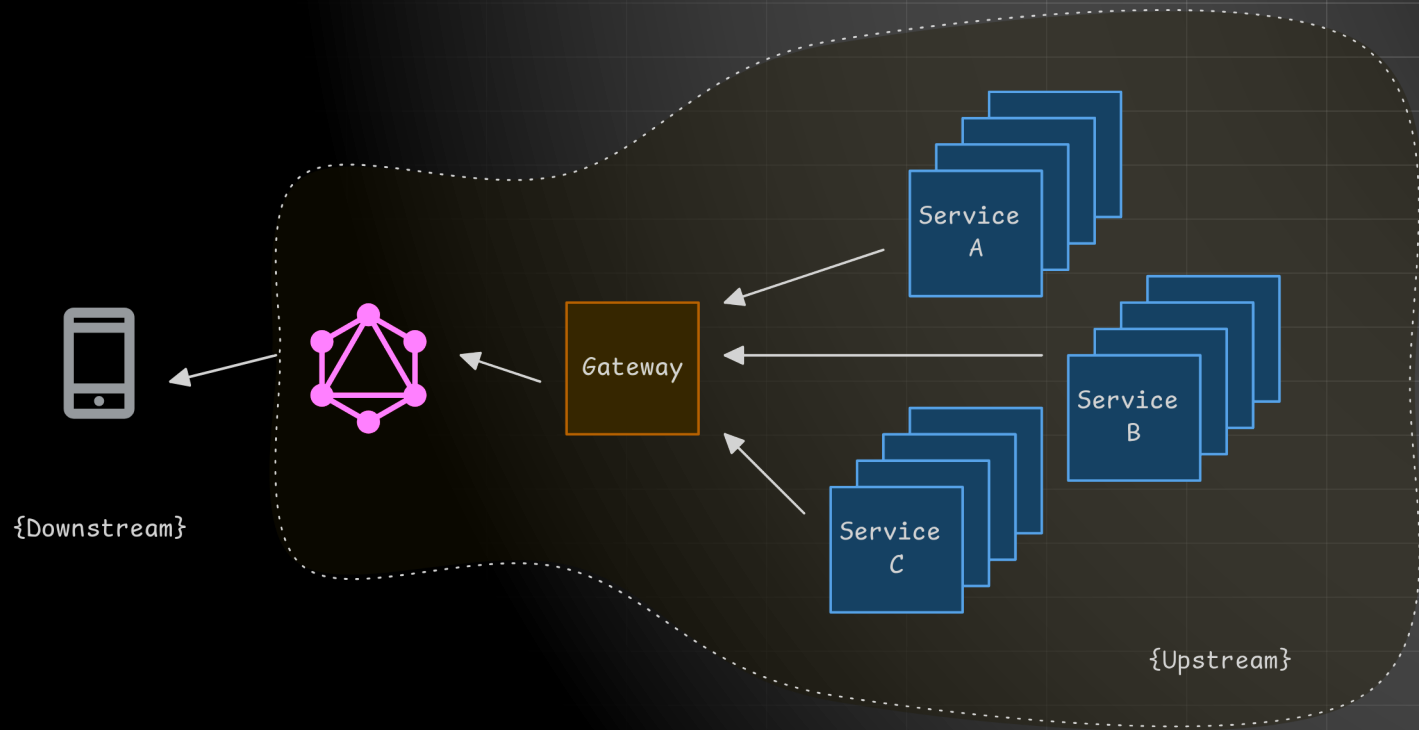


GraphQL?

Architecture



Architecture



GraphQL as client-side abstraction

2022 Dream11

500

ENGINEERS

150

MICROSERVICES

10,000,000

CONCURRENCY

200,000,000

USERS

2022 Dream11

500

ENGINEERS

150

MICROSERVICES

10,000,000

CONCURRENCY

200,000,000

USERS



GraphQL

Saved the day...



Not so fast!



Part 2

Challenges

1. Performance

1. Performance

👍 Latency & Throughput

1. Performance

👍 Latency & Throughput

👎 Infrastructure

1. Performance

👍 Latency & Throughput

👎 Infrastructure

50,000

CORES

Things that worked for us

Things that worked for us

👉 Micro Optimizations

Things that worked for us

👉 Micro Optimizations

👉 Infrastructure Tuning

Things that worked for us

- 👉 Micro Optimizations
- 👉 Infrastructure Tuning
- 👉 Benchmarking

Things that worked for us

- 👉 Micro Optimizations
- 👉 Infrastructure Tuning
- 👉 Benchmarking
- 👉 Caching on Upstream

Things that worked for us

- 👉 Micro Optimizations
- 👉 Infrastructure Tuning
- 👉 Benchmarking
- 👉 Caching on Upstream

6,400

CORES

Things that worked for us

- 👉 Micro Optimizations
- 👉 Infrastructure Tuning
- 👉 Benchmarking
- 👉 Caching on Upstream

6,400

CORES



High Priority

2. Reliability

2. Reliability

 Fragility

2. Reliability

 Fragility

2. Reliability

 Fragility

 Resiliency

Things that worked for us

Things that worked for us

👉 Redefining Processes

Things that worked for us

👉 Redefining Processes

👉 Micro-optimizing Cold Paths

Things that worked for us

- 👉 Redefining Processes
- 👉 Micro-optimizing Cold Paths
- 👉 Rate Limiting & Circuit Breaking

Things that worked for us

- 👉 Redefining Processes
- 👉 Micro-optimizing Cold Paths
- 👉 Rate Limiting & Circuit Breaking

✗ Query Cost

3. Maintainability

3. Maintainability

- ? Library Upgrades
- ? Unused Nodes
- ? Data Loaders
- ? Unimplemented Resolver
- ? Code Duplication
- ? ...

3. Maintainability

- ? Library Upgrades
- ? Unused Nodes
- ? Data Loaders
- ? Unimplemented Resolver
- ? Code Duplication
- ? ...



Summary

1. Performance



2. Reliability



3. Maintenance



Part 3

Learnings of 8 Years

 **Liberties Constraint and Constraints**
Liberate 

Liberties



ViralHog

Constraints



GraphQL

1. Schema

2. Query

3. Resolver

Resolver

Conformance requirements expressed as algorithms can be fulfilled by an implementation of this specification **in any way as long as the perceived result is equivalent**. Algorithms described in this document are written to be easy to understand. Implementers are encouraged to include equivalent but optimized implementations.

Step 1



Avoid Business Logic



Allow only Orchestration Logic

Step 2



No Handwritten Resolver



Configuration Based

Step 3

 Learn from SQL



Part 4

Generalized Runtime for GraphQL

`>_ tailcall`

Built with ❤️ using
Apache 2.0 License
High Performance



How does it work?

GraphQL Schema 🙋

```
schema {
  query: Query
}

type Query {
  posts: [Post]
}

type Post {
  id: ID!
  title: String!
  body: String
  userId: ID!
  user: User
}

type User {
  id: ID!
  name: String!
  email: String!
}
```

Annotate

```
schema @upstream(baseUrl: "https://api.d11.local") {
  query: Query
}

type Query {
  posts: [Post] @http(path: "/posts")
}

type Post {
  id: ID!
  title: String!
  body: String
  userId: ID!
  user: User @http(path: "/users/{value.userId}")
}

type User {
  id: ID!
  name: String!
  email: String!
}
```

Start the server

```
> tailcall start config.graphql  
INFO File read: config.graphql ... ok  
INFO 🚀 Tailcall launched at [127.0.0.1:8000] over HTTP/1.1
```

First Class Support

- REST
- gRPC
- GraphQL

First Class Support

- REST
- gRPC
- GraphQL

@addField
@call
@graphql
@http
@modify
@rest
@telemetry

@cache
@expr
@grpc
@link
@omit
@server
@upstream

First Class Support

- REST
- gRPC
- GraphQL

```
@addField      @cache
               @call      @expr
               @graphql   @grpc
               @http      @link
               @modify    @omit
               @rest      @server
               @telemetry @upstream
```



Router

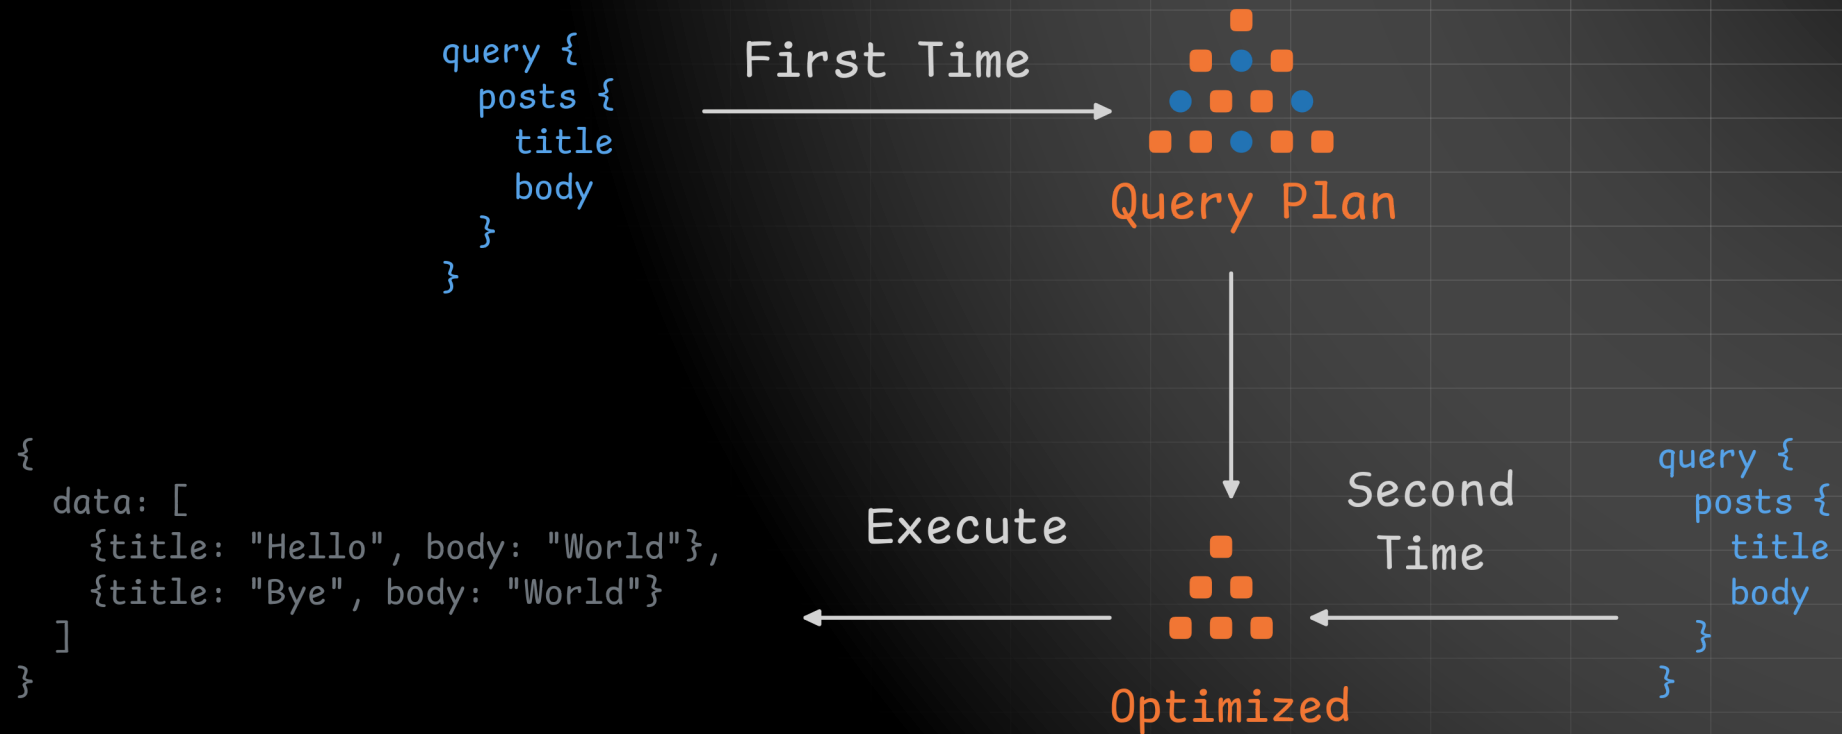
Constraints that Liberate

2 Key Capabilities

AOT Analyzer

- › `tailcall` start config.graphql
INFO File read: config.graphql ... ok
INFO N + 1 detected: 1
ERROR Invalid Configuration
Caused by:
 - argument 'id' is a nullable type [at Query.args.id]

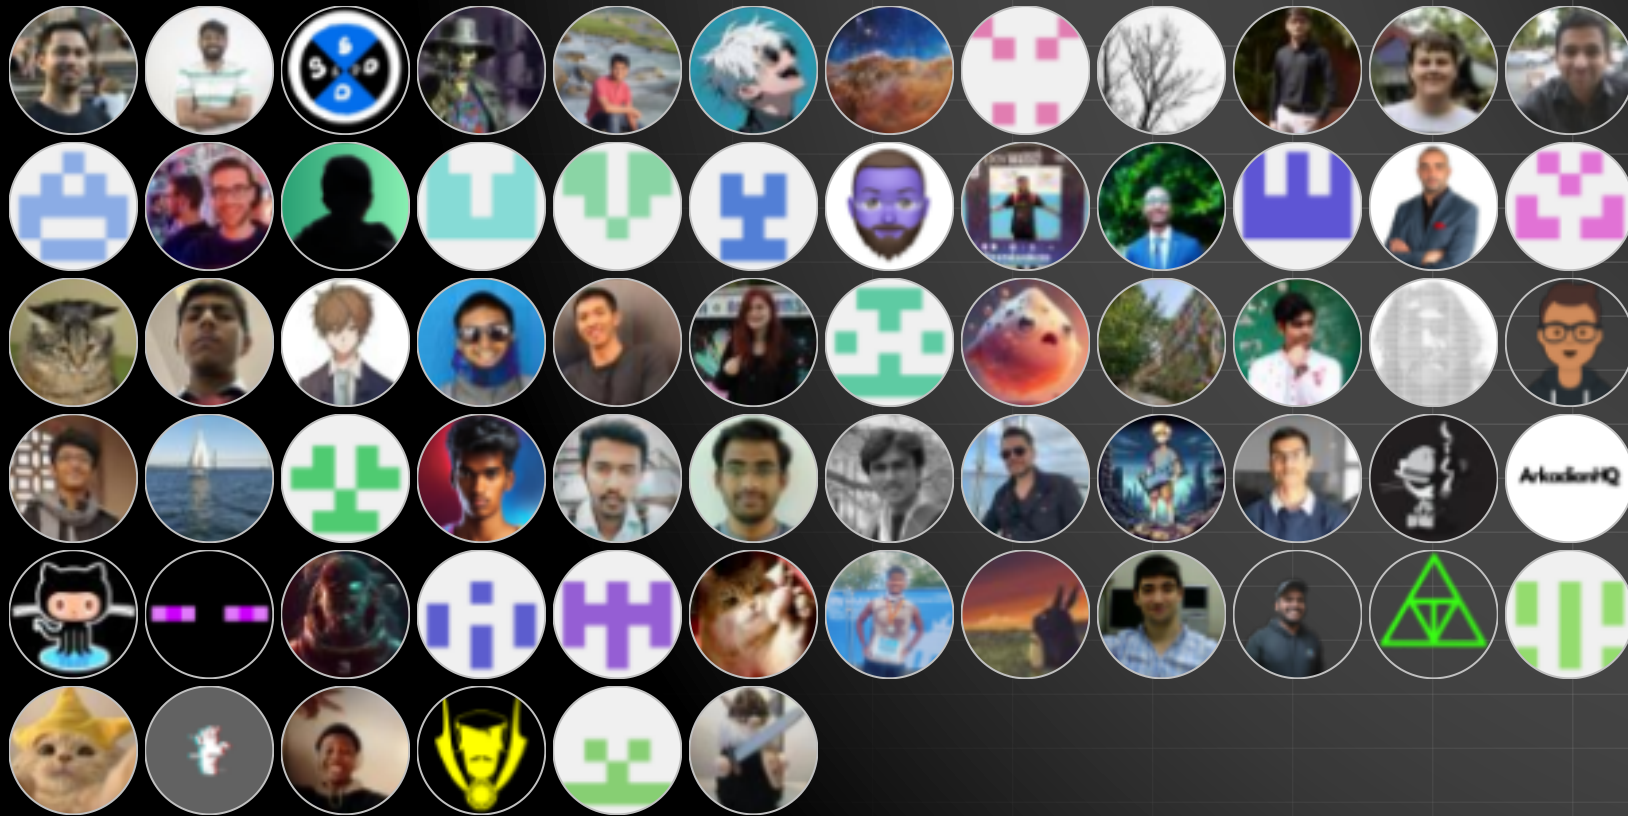
JIT Optimizer



Revisiting the Challenges

	Before	<code>>_</code> tailcall
1. Performance	★ ★ ☆	★ ★ ★
2. Reliability	★ ★ ☆	★ ★ ★
3. Maintenance	☆ ☆ ☆	★ ★ ★
4. <u>Flexibility</u>	★ ★ ★	★ ☆ ☆

Awesome Community 🙌



Part 5

Takeaway

Takeaway

1. Innovation on GraphQL performance is necessary.
2. Handwritten GraphQL is difficult to maintain.
3. Library authors should take inspiration from SQL engines.

#TailcallHack

- Build the fastest GraphQL Server
- Checkout - [tailcallhq/hackathon](https://tailcallhq.com/hackathon)

\$ 5,000



Thank You! ❤️

Tushar Mathur

1. Founder of Tailcall
2. Ex VP of Engineering at Dream11 (2016-2022)

 @tusharmath

 @tusharmath

 @tusharmath